



IIT KHARAGPUR



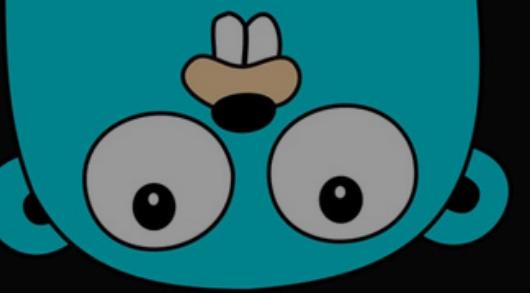
TITLE:

GO CONCURRENCY



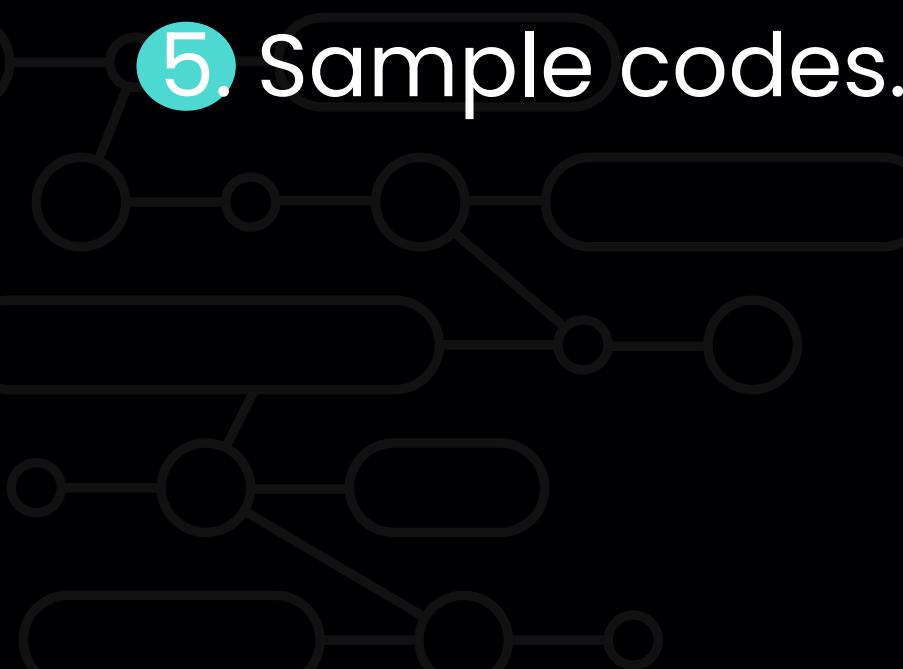
~ADITYA DEBNATH





CONTENTS :

1. Brief intro of GoLanguage.
2. Concurrent VS Sequential Programing
3. Benefits of Concurrent Programming
4. How concurrency is implemented in GO (using Go routines, channels, waitgroups and mutexes)
5. Sample codes.



GO LANGUAGE

Developed by
Google, 2007



Free and open
source language



Compiled
language



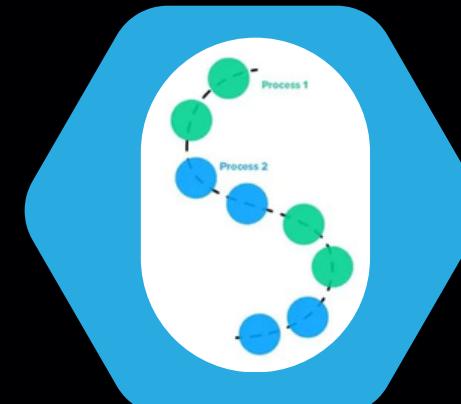
Fast compilation
and execution



Easy in syntax



Supports
concurrency



Sequential V/S Concurrent Programming

1 Single Execution Path

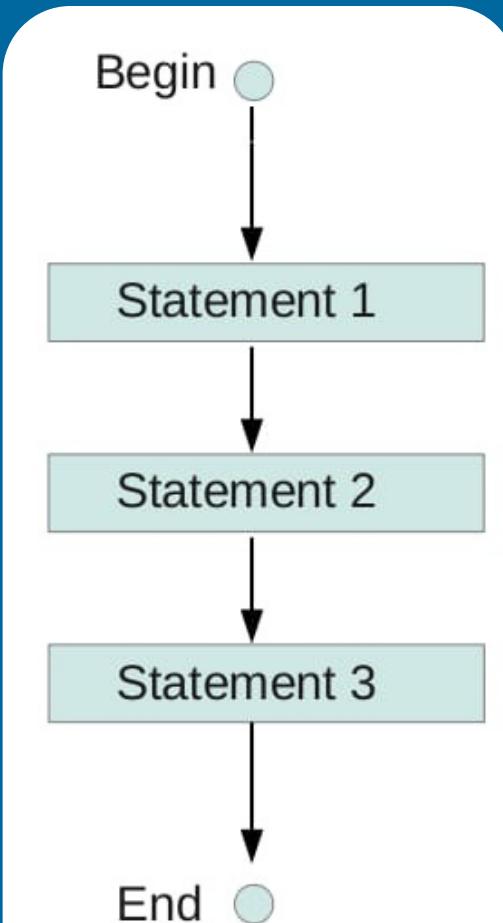
2 Deterministic Execution

3 Synchronous Execution

4 Limited Scalability

5 Readability

6 Ex. coding in C or C++



1 Parallel Execution

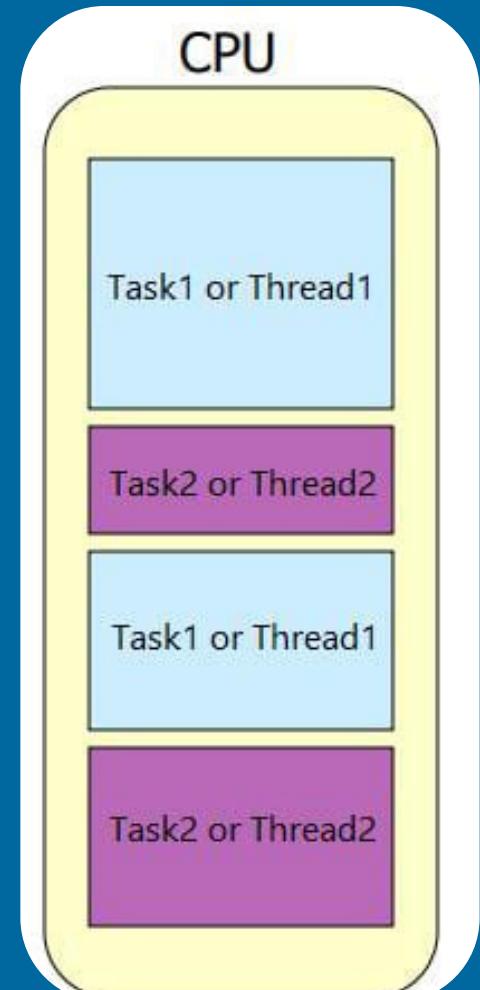
2 Not always deterministic

3 Asynchronous Programming

4 Optimally scalable

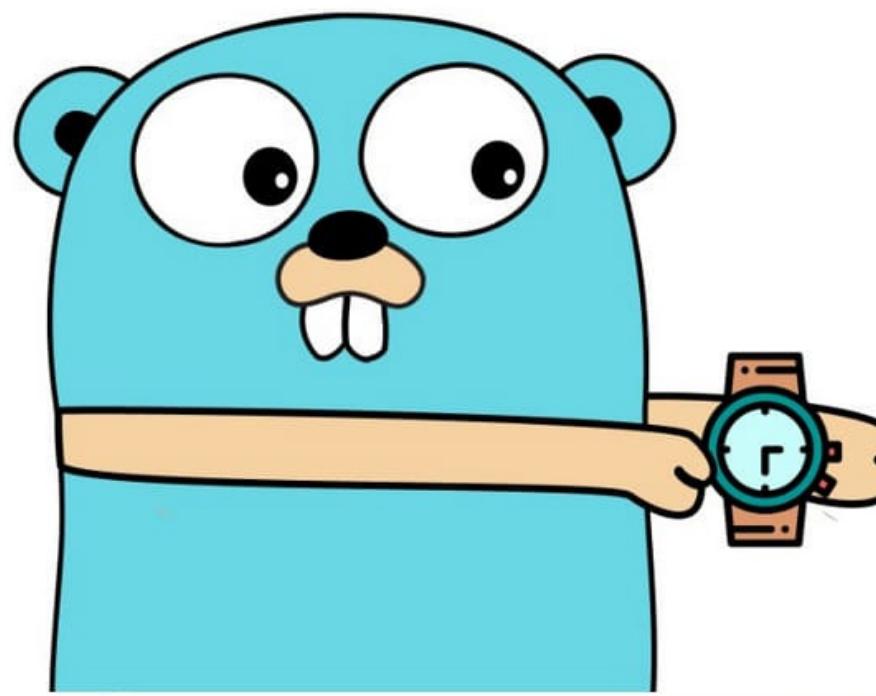
5 Not easily readable

6 Ex. go routines func, in go



BENEFITS OF CONCURRENT PROGRAMMING

It's Go Time



1 IMPROVED PERFORMANCE

2 Enhanced Responsiveness

3 Resource Utilization

4 Simplified Complex Systems:

5 Parallel Algorithms

How it's done then?

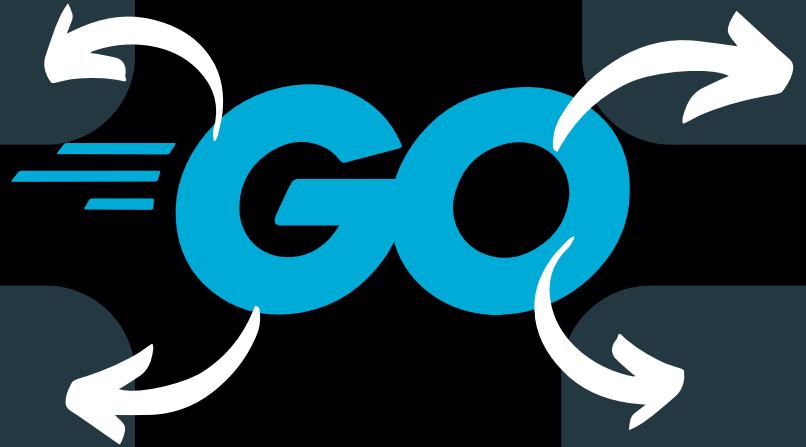
Go provides various fundamental blocks that help make concurrent programming user-friendly.

GO ROUTINES

CHANNELS

WAITGROUPS

MUTEXES





GO ROUTINES

- Goroutines are functions that run concurrently with other functions. They are lightweight threads managed by the Go runtime, allowing thousands or even millions of them to be created within a single Go program without significant overhead.
- Goroutines are created using the `go` keyword followed by a function call. When a function is called with `go`, it starts executing concurrently in its own goroutine, allowing the program to proceed with other tasks without waiting for it to finish.
- Goroutines are multiplexed onto a smaller number of OS threads, meaning that many goroutines can run on a single operating system thread, which makes them very efficient.
- Goroutines are a key feature of Go's approach to concurrency, enabling developers to write highly concurrent and efficient programs with ease.

```
1 package main
2
3 import (
4     "fmt"
5     "time"
6 )
7
8 func sayHello() {
9     fmt.Println("Hello from
10    goroutine!")
11 }
12 func main() {
13     // Start a new goroutine
14     go sayHello()
15
16     // Print from the main goroutine
17     fmt.Println("Hello from main!")
18
19     // Add a delay to allow the
20     // goroutine to execute
21     time.Sleep(100 * time
22         .Millisecond)
23 }
```

Run

```
go run /tmp/2JNAnGORjT.go
Hello from main!
Hello from goroutine!
```

CHANNELS

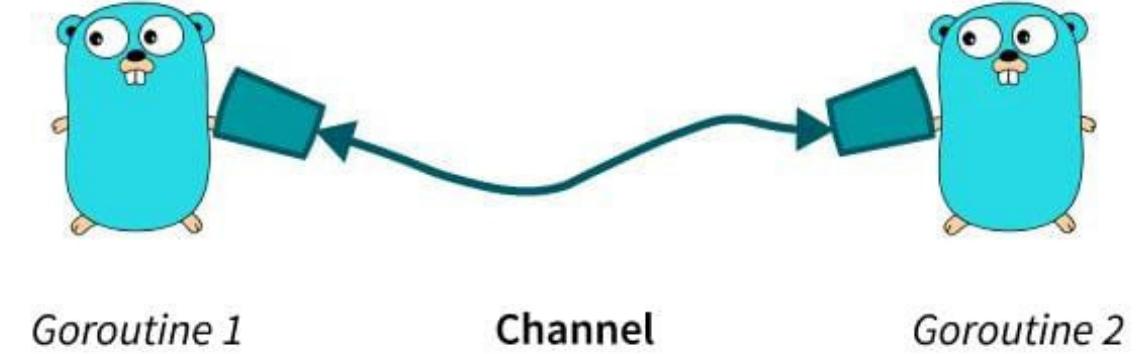
Channels are the pipes that connect concurrent goroutines. You can send values into channels from one goroutine and receive those values into another goroutine.

Create a new channel with `make(chan val-type)`. Channels are typed by the values they convey.

Send a value into a channel using the `channel <-` syntax. Here we send "ping" to the messages channel we made above, from a new goroutine.

The `<-channel` syntax receives a value from the channel. Here we'll receive the "ping" message we sent above and print it out.

By default sends and receives block until both the sender and receiver are ready. This property allowed us to wait at the end of our program for the "ping" message without having to use any other synchronization



```
package main  
  
import "fmt"  
  
func main() {  
  
    messages := make(chan  
        string)  
  
    go func() { messages <-\n        "ping" }()  
  
    msg := <-messages  
    fmt.Println(msg)  
}
```

~ BUFFERED CHANNELS

By default channels are *unbuffered*, meaning that they will only accept sends (`chan <-`) if there is a corresponding receive (`<- chan`) ready to receive the sent value. *Buffered channels* accept a limited number of values without a corresponding receiver for those values.

Here we make a channel of strings buffering up to 2 values

Because this channel is buffered, we can send these values into the channel without a corresponding concurrent receive.

package main

import "fmt"

func main() {

messages := make(chan string, 2)

messages <- "buffered"

messages <- "channel"

fmt.Println(<-messages)

fmt.Println(<-messages)

}

WAITGROUPS

This WaitGroup is used to wait for all the goroutines launched here to finish. Note: if a WaitGroup is explicitly passed into functions, it should be done by pointer.

Launch several goroutines and increment the WaitGroup counter for each.

Block until the WaitGroup counter goes back to 0; all the workers notified they're done.

```
package main

import (
    "fmt"
    "sync"
    "time"
)

func worker(id int) {
    fmt.Printf("Worker %d starting\n", id)

    time.Sleep(time.Second)
    fmt.Printf("Worker %d done\n", id)
}

func main() {

    var wg sync.WaitGroup

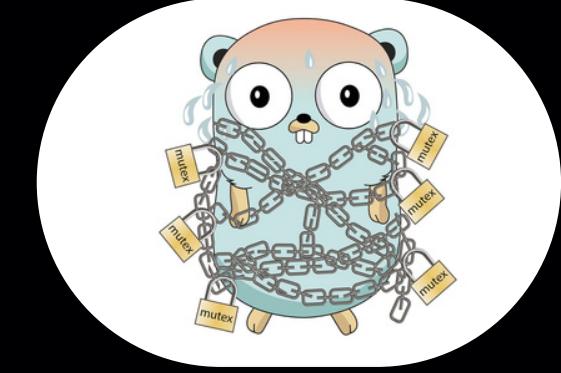
    for i := 1; i <= 5; i++ {
        wg.Add(1)

        go func() {
            defer wg.Done()
            worker(i)
        }()
    }

    wg.Wait()

}
```

MUTEXES



In Go, a mutex (short for mutual exclusion) is a synchronization primitive used to protect shared resources from being accessed concurrently by multiple goroutines. Mutexes ensure that only one goroutine can access the shared resource at any given time, preventing race conditions and data races.

Mutexes work by locking and unlocking. When a goroutine wants to access a shared resource, it locks the mutex. If the mutex is already locked by another goroutine, the goroutine will be blocked until the mutex becomes available. Once the goroutine finishes its work with the shared resource, it unlocks the mutex to allow other goroutines to access it.

```
1 package main
2
3 import (
4     "fmt"
5     "sync"
6     "time"
7 )
8
9 var (
10    counter int
11    mutex   sync.Mutex
12 )
13
14 func increment() {
15     mutex.Lock()
16     defer mutex.Unlock()
17     counter++
18     fmt.Println("Incrementing counter to", counter)
19 }
20
21 func main() {
22     for i := 0; i < 5; i++ {
23         go increment()
24     }
25
26     time.Sleep(time.Second) // Waiting for goroutines to finish
27     fmt.Println("Final counter value:", counter)
28 }
29
```

[Run](#)

//code time </>

```
1 package main
2
3 import (
4     "fmt"
5     "sync"
6 )
7
8 func worker(id int, jobs <-chan int,
9             results chan<- int, wg *sync.
10            WaitGroup, mu *sync.Mutex) {
11    defer wg.Done()
12    for j := range jobs {
13        fmt.Printf("Worker %d
14             processing job %d\n",
15             id, j)
16        // Simulating some work
17        // Protecting shared
18        // resource with mutex
19        mu.Lock()
20        results <- j * 2
21        mu.Unlock()
22    }
23
24    func main() {
25        // Number of workers
26        numWorkers := 3
27        // Number of jobs
28        numJobs := 5
29
30        // Create channels
31        jobs := make(chan int, numJobs)
32        results := make(chan int,
33                      numJobs)
```

```
34        // Create wait group and mutex
35        var wg sync.WaitGroup
36        var mu sync.Mutex
37
38        // Add workers to wait group
39        for i := 1; i <= numWorkers; i++ {
40            wg.Add(1)
41            go worker(i, jobs, results,
42                       &wg, &mu)
43        }
44
45        // Send jobs to workers
46        for j := 1; j <= numJobs; j++ {
47            jobs <- j
48        }
49        close(jobs)
50
51        // Wait for all workers to
52        // finish
53        wg.Wait()
54
55        // Close results channel after
56        // all workers are done
57        close(results)
58
59        // Collect and print results
60        for r := range results {
61            fmt.Println("Result:", r)
62        }
63    }
64 }
```

main.go

Output

go run /tmp/d18cMumM90.go

Worker 3 processing job 1

Worker 3 processing job 2

Worker 3 processing job 3

Worker 3 processing job 4

Worker 1 processing job 5

Result: 2

Result: 4

Result: 6

Result: 8

Result: 10

THANK YOU

